

CS5720

Neural Networks & Deep Learning - ICP-5

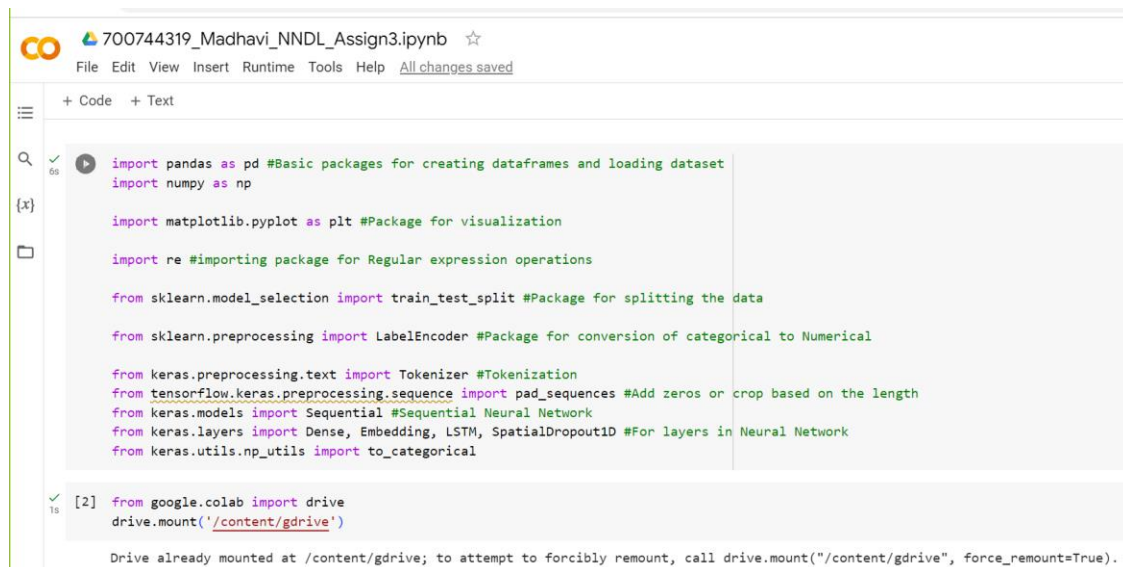
Name: Madhavi Mamidala

#700: 700744319

Username: MXM43190

GitHub Link:

https://github.com/MadhaviMamidala/CS5720_ASSIGNMENT5_700744319/blob/main/700744319_Madhavi_NNDL_Assign5.ipynb



The screenshot shows a Jupyter Notebook titled "700744319_Madhavi_NNDL_Assign3.ipynb". The interface includes a top menu bar with options like File, Edit, View, Insert, Runtime, Tools, and Help. Below the menu, there are tabs for "+ Code" and "+ Text". The main area displays two code cells. The first cell contains a series of import statements for various libraries, each followed by a comment explaining its purpose. The second cell contains a single line of code to mount Google Drive, followed by a message indicating that the drive is already mounted.

```
import pandas as pd #Basic packages for creating dataframes and loading dataset
import numpy as np

import matplotlib.pyplot as plt #Package for visualization

import re #importing package for Regular expression operations

from sklearn.model_selection import train_test_split #Package for splitting the data

from sklearn.preprocessing import LabelEncoder #Package for conversion of categorical to Numerical

from keras.preprocessing.text import Tokenizer #Tokenization
from tensorflow.keras.preprocessing.sequence import pad_sequences #Add zeros or crop based on the length
from keras.models import Sequential #Sequential Neural Network
from keras.layers import Dense, Embedding, LSTM, SpatialDropout1D #For layers in Neural Network
from keras.utils.np_utils import to_categorical

[2]: from google.colab import drive
    drive.mount('/content/gdrive')
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).

{x}

✓
0s

```
[3] import pandas as pd

path_to_csv = '/content/gdrive/MyDrive/Sentiment.csv'
# Load the dataset as a Pandas DataFrame
dataset = pd.read_csv(path_to_csv, header=0)

# Select only the necessary columns 'text' and 'sentiment'
mask = dataset.columns.isin(['text', 'sentiment'])
data = dataset.loc[:, mask]

# Keeping only the necessary columns
```

✓
0s

cleaning text , removing all special characters

```
data['text'] = data['text'].apply(lambda x: x.lower())
data['text'] = data['text'].apply(lambda x: re.sub('[^a-zA-Z0-9\s]', '', x))
```

<ipython-input-4-8bc15de6ba7d>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
data['text'] = data['text'].apply(lambda x: x.lower())
<ipython-input-4-8bc15de6ba7d>:4: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
data['text'] = data['text'].apply(lambda x: re.sub('[^a-zA-Z0-9\s]', '', x))
```

✓
3s

```
[5] for idx, row in data.iterrows():
    row[0] = row[0].replace('rt', ' ') #Removing Retweets
```

✓
1s

```
[6] max_fatures = 2000
tokenizer = Tokenizer(num_words=max_fatures, split=' ') #Maximum words is 2000 to tokenize sentence
tokenizer.fit_on_texts(data['text'].values)
X = tokenizer.texts_to_sequences(data['text'].values) #taking values to feature matrix
```

{x}

✓
0s

```
[7] X = pad_sequences(X) #Padding the feature matrix

embed_dim = 128 #Dimension of the Embedded layer
lstm_out = 196 #Long short-term memory (LSTM) layer neurons
```

✓
0s

Sequential Model Creation

```
def createmodel():
    model = Sequential() #Sequential Neural Network
    model.add(Embedding(max_fatures, embed_dim,input_length = X.shape[1])) #input dimension 2000 Neurons, output dimension 128 Neurons
    model.add(LSTM(lstm_out, dropout=0.2, recurrent_dropout=0.2)) #Drop out 20%, 196 output Neurons, recurrent dropout 20%
    model.add(Dense(3,activation='softmax')) #3 output neurons[positive, Neutral, Negative], softmax as activation
    model.compile(loss = 'categorical_crossentropy', optimizer='adam',metrics = ['accuracy']) #Compiling the model
    return model
# print(model.summary())
```

✓
0s

```
[9] # Label Encoding of values
labelencoder = LabelEncoder() #Applying label Encoding on the label matrix
integer_encoded = labelencoder.fit_transform(data['sentiment']) #fitting the model
y = to_categorical(integer_encoded)
X_train, X_test, Y_train, Y_test = train_test_split(X,y, test_size = 0.33, random_state = 42) #67% training data, 33% test data split
```

```
[10] # Model Training
batch_size = 32 #Batch size 32
model = createmodel() #Function call to Sequential Neural Network
model.fit(X_train, Y_train, epochs = 1, batch_size=batch_size, verbose = 2) #verbose the higher, the more messages
score,acc = model.evaluate(X_test,Y_test,verbose=2,batch_size=batch_size) #evaluating the model
print(score)
print(acc)

WARNING:tensorflow:Layer lstm will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
291/291 - 59s - loss: 0.8294 - accuracy: 0.6450 - 59s/epoch - 202ms/step
144/144 - 2s - loss: 0.7528 - accuracy: 0.6710 - 2s/epoch - 12ms/step
0.7528498768806458
0.6710354089736938

[11] print(model.metrics_names) #metrics of the model

['loss', 'accuracy']
```

1. Save the model and use the saved model to predict on new text data (ex, "A lot of good things are happening. We are respected again throughout the world, and that's a great thing.@realDonaldTrump")

```
model.save('sentimentAnalysis.h5') #Saving the model

from keras.models import load_model #Importing the package for importing the saved model
model= load_model('sentimentAnalysis.h5') #loading the saved model

WARNING:tensorflow:Layer lstm will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
```

```
[14] print(integer_encoded)
print(data['sentiment'])

[ 1  2  1 ...  2  0  2]
0      Neutral
1      Positive
2      Neutral
3      Positive
4      Positive
...
13866   Negative
13867   Positive
13868   Positive
13869   Negative
13870   Positive
Name: sentiment, Length: 13871, dtype: object
```

{x} ✓ 06



```
# Predicting on the text data
sentence = ['A lot of good things are happening. We are respected again throughout the world, and that is a great thing.@realDonaldTrump']

sentence = tokenizer.texts_to_sequences(sentence) # Tokenizing the sentence
sentence = pad_sequences(sentence, maxlen=28, dtype='int32', value=0) # Padding the sentence
sentiment_probs = model.predict(sentence, batch_size=1, verbose=2)[0] # Predicting the sentence text
sentiment = np.argmax(sentiment_probs)

print(sentiment_probs)
if sentiment == 0:
    print("Neutral")
elif sentiment < 0:
    print("Negative")
elif sentiment > 0:
    print("Positive")
else:
    print("Cannot be determined")
```

```
1/1 - 0s - 370ms/epoch - 370ms/step
[0.733563  0.12827614 0.13816082]
Neutral
```

{x}



```
from keras.wrappers.scikit_learn import KerasClassifier #importing Keras classifier
from sklearn.model_selection import GridSearchCV #importing Grid search CV

model = KerasClassifier(build_fn=createmodel,verbose=2) #initiating model to test performance by applying multiple hyper parameters
batch_size = [10, 20, 40] #hyper parameter batch_size
epochs = [1, 2] #hyper parameter no. of epochs
param_grid = {'batch_size':batch_size, 'epochs':epochs} #creating dictionary for batch size, no. of epochs
grid = GridSearchCV(estimator=model, param_grid=param_grid) #Applying dictionary with hyper parameters
grid_result = grid.fit(X_train,Y_train) #Fitting the model
# summarize results
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_)) #Best score, best hyper parameters

... <ipython-input-16-6c99b49150f4>:4: DeprecationWarning: KerasClassifier is deprecated, use Sci-Keras (https://github.com/adriangb/scikeras) instead. See https://www.adriangb.com/
model = KerasClassifier(build_fn=createmodel,verbose=2) #initiating model to test performance by applying multiple hyper parameters
WARNING:tensorflow:Layer lstm_1 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
744/744 - 115s - loss: 0.8252 - accuracy: 0.6438 - 115s/epoch - 155ms/step
186/186 - 3s - loss: 0.7521 - accuracy: 0.6622 - 3s/epoch - 16ms/step
WARNING:tensorflow:Layer lstm_2 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
744/744 - 113s - loss: 0.8207 - accuracy: 0.6493 - 113s/epoch - 151ms/step
186/186 - 3s - loss: 0.7599 - accuracy: 0.6746 - 3s/epoch - 16ms/step
WARNING:tensorflow:Layer lstm_3 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
744/744 - 112s - loss: 0.8261 - accuracy: 0.6438 - 112s/epoch - 151ms/step
186/186 - 2s - loss: 0.7709 - accuracy: 0.6751 - 2s/epoch - 12ms/step
WARNING:tensorflow:Layer lstm_4 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
744/744 - 111s - loss: 0.8249 - accuracy: 0.6469 - 111s/epoch - 149ms/step
186/186 - 2s - loss: 0.7637 - accuracy: 0.6738 - 2s/epoch - 12ms/step
```

```
186/186 - 3s - loss: 0.7521 - accuracy: 0.6622 - 3s/epoch - 16ms/step
WARNING:tensorflow:Layer lstm_2 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
744/744 - 113s - loss: 0.8207 - accuracy: 0.6493 - 113s/epoch - 151ms/step
***
186/186 - 3s - loss: 0.7599 - accuracy: 0.6746 - 3s/epoch - 16ms/step
WARNING:tensorflow:Layer lstm_3 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
744/744 - 112s - loss: 0.8261 - accuracy: 0.6438 - 112s/epoch - 151ms/step
186/186 - 2s - loss: 0.7709 - accuracy: 0.6751 - 2s/epoch - 12ms/step
WARNING:tensorflow:Layer lstm_4 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
186/186 - 2s - loss: 0.8249 - accuracy: 0.6469 - 111s/epoch - 149ms/step
186/186 - 2s - loss: 0.7637 - accuracy: 0.6738 - 2s/epoch - 12ms/step
WARNING:tensorflow:Layer lstm_5 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
744/744 - 124s - loss: 0.8216 - accuracy: 0.6433 - 124s/epoch - 166ms/step
186/186 - 2s - loss: 0.7893 - accuracy: 0.6464 - 2s/epoch - 12ms/step
WARNING:tensorflow:Layer lstm_6 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
Epoch 1/2
744/744 - 115s - loss: 0.8195 - accuracy: 0.6500 - 115s/epoch - 154ms/step
Epoch 2/2
744/744 - 100s - loss: 0.6742 - accuracy: 0.7154 - 100s/epoch - 134ms/step
186/186 - 2s - loss: 0.7385 - accuracy: 0.6789 - 2s/epoch - 12ms/step
WARNING:tensorflow:Layer lstm_7 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
Epoch 1/2
744/744 - 114s - loss: 0.8251 - accuracy: 0.6469 - 114s/epoch - 153ms/step
Epoch 2/2
744/744 - 101s - loss: 0.6831 - accuracy: 0.7113 - 101s/epoch - 136ms/step
186/186 - 3s - loss: 0.7479 - accuracy: 0.6622 - 3s/epoch - 14ms/step
WARNING:tensorflow:Layer lstm_8 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
Epoch 1/2
744/744 - 115s - loss: 0.8268 - accuracy: 0.6457 - 115s/epoch - 155ms/step
Epoch 2/2
744/744 - 102s - loss: 0.6753 - accuracy: 0.7146 - 102s/epoch - 138ms/step
186/186 - 3s - loss: 0.7547 - accuracy: 0.6772 - 3s/epoch - 18ms/step
WARNING:tensorflow:Layer lstm_9 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
Epoch 1/2
744/744 - 114s - loss: 0.8222 - accuracy: 0.6451 - 114s/epoch - 154ms/step
Epoch 2/2
744/744 - 105s - loss: 0.6719 - accuracy: 0.7116 - 105s/epoch - 141ms/step
186/186 - 3s - loss: 0.7599 - accuracy: 0.6728 - 3s/epoch - 17ms/step
WARNING:tensorflow:Layer lstm_10 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
Epoch 1/2
```